

Aviva x AWS

How Aviva deployed enterprise MLOps to automate machine learning and drive £3 billion in annual claims efficiency.



Aviva at a Glance

Global Reach & Heritage

- Founded 1696—one of the world's oldest **financial services** organisations
- Operating across 16 countries
- Serving 33+ million customers worldwide

Core Offerings

- Life and general insurance
- Pensions and investment services
- **70+ active AI and ML** use cases in production

Business Problems



Time-Consuming Operational Tasks

Data scientists spend over **50% of their time** on operational tasks, diverting focus from innovative model development.



High On-Premise Infrastructure Costs

Inefficient resource utilization and **high maintenance expenses**, driving significant fixed costs.



Slow, Error-Prone Deployment

An estimated **50% of ML projects never reach production**, indicating significant deployment hurdles and wasted investment.



Massive Claims Volume & Cost

Aviva processes **400,000 annual claims**, amounting to **3 billion EUR**, demanding enterprise-grade automation and governance for optimized workflows.

The solution:

AWS Enterprise MLOps Framework

Implement a modern, **cloud-native MLOps architecture** built on **four building blocks** designed for enterprise scale, security, and governance.

4 Building Blocks

1. Network Infrastructure

- 3 isolated AWS accounts: Development, Staging and Production
- new network infrastructure set up across all environments via VPCs and endpoints; zero-trust security and compliance built in.

2. Amazon SageMaker Studio

- Unified workspace for ML development, training, and deployment
- native collaboration, versioning, and monitoring.

3. Amazon Project Templates

- Automated ML pipelines enforcing best practices, CI/CD, security across environments for consistency at scale.

4. Seed Code

- Production-ready scaffolding, preprocessing scripts, and baseline pipeline code
- accelerating experimentation and standardising patterns



The Remedy Use Case: Complex Orchestration of 14 ML Models

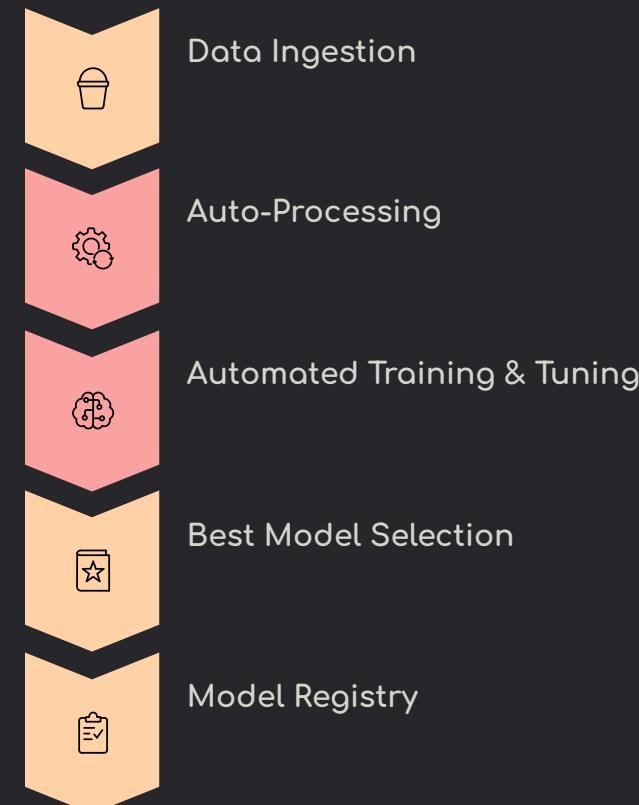
- The Remedy process involves deciding **whether to repair or write off a damaged vehicle** — a **key decision** in Aviva's operations.
- Managing **14 different machine learning models** and business rules on outdated legacy systems caused **inefficiencies**.

Transforming Model Development

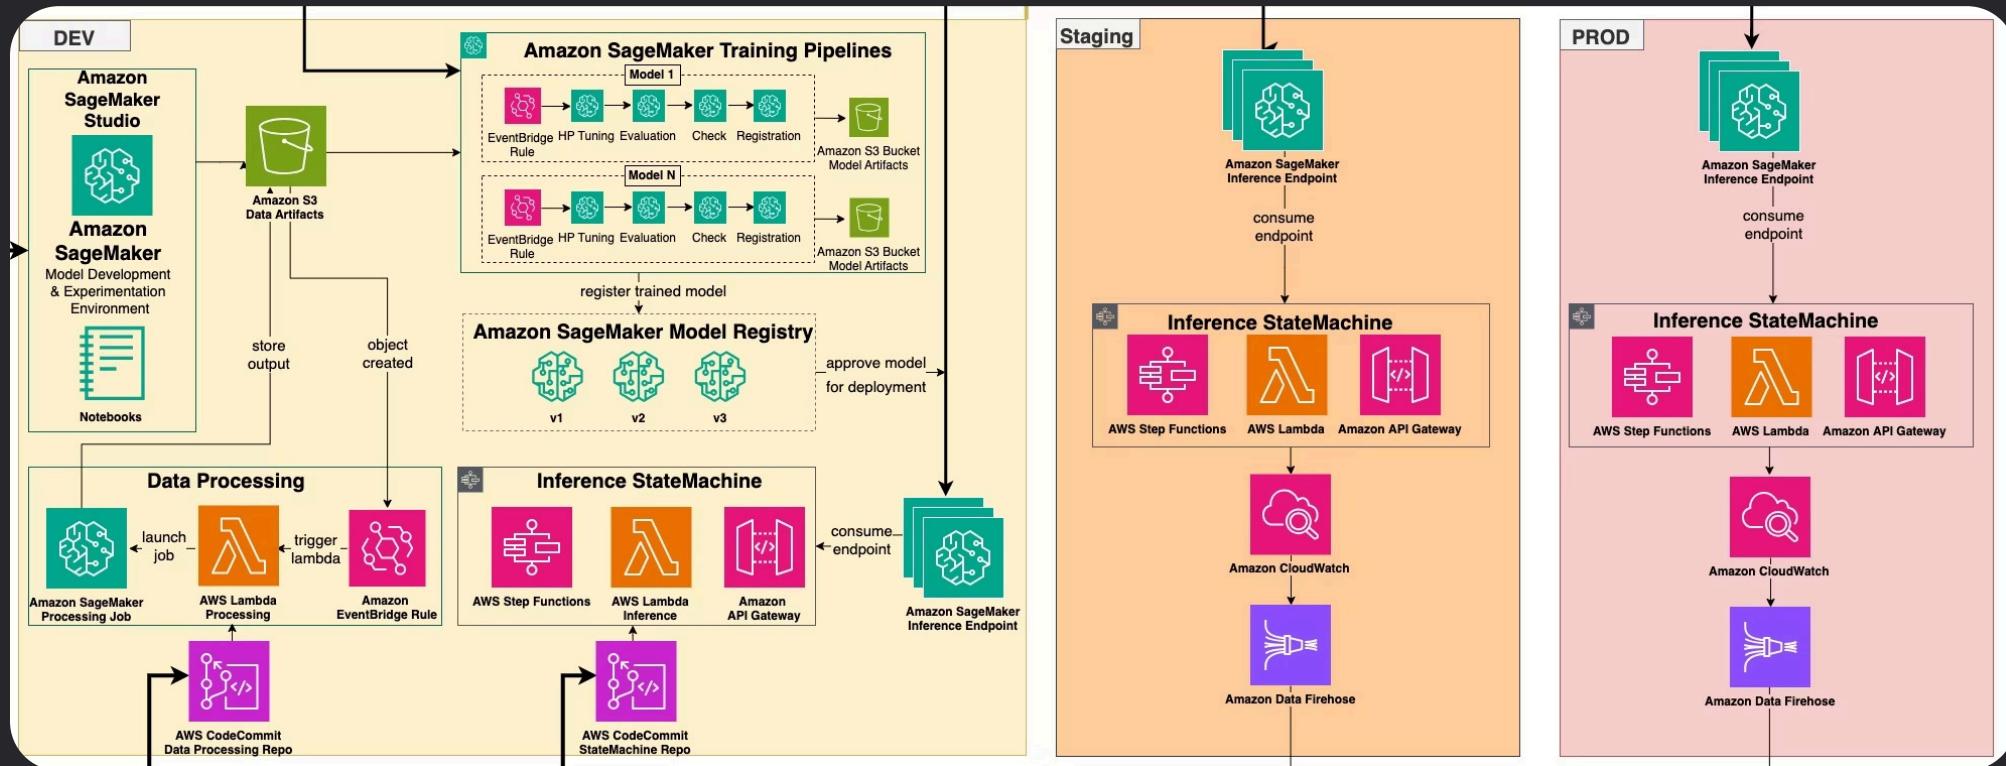
Manual Development



Automated Development Pipeline



Overview MLOps Architecture



Development



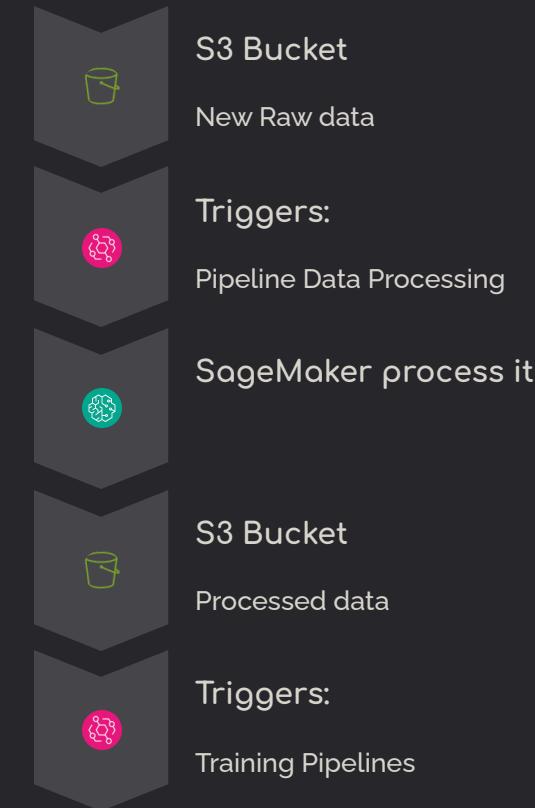
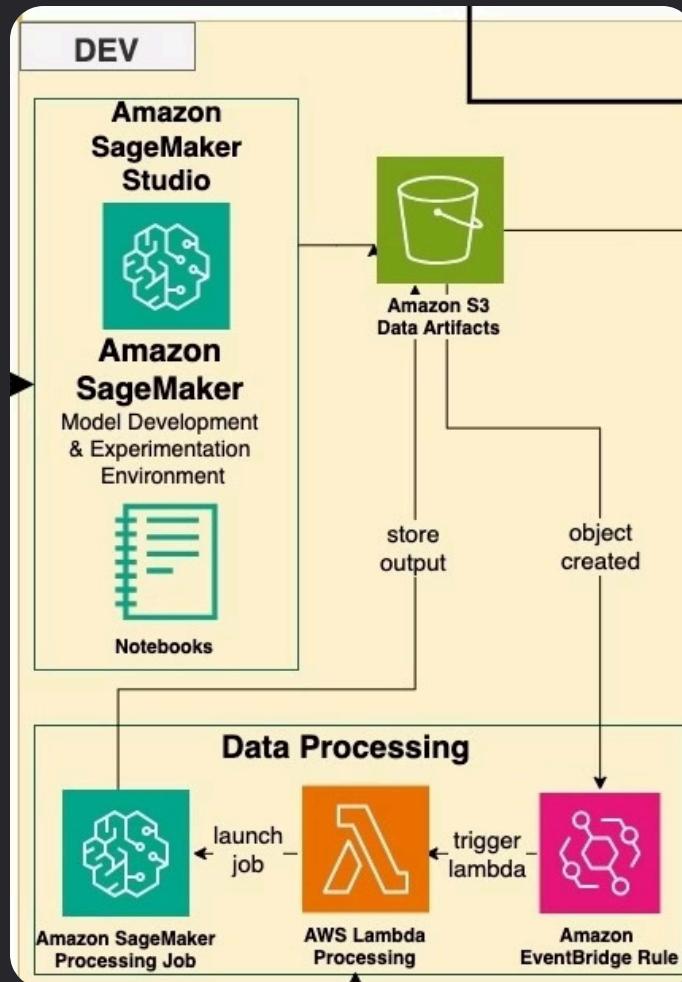
Staging



Production

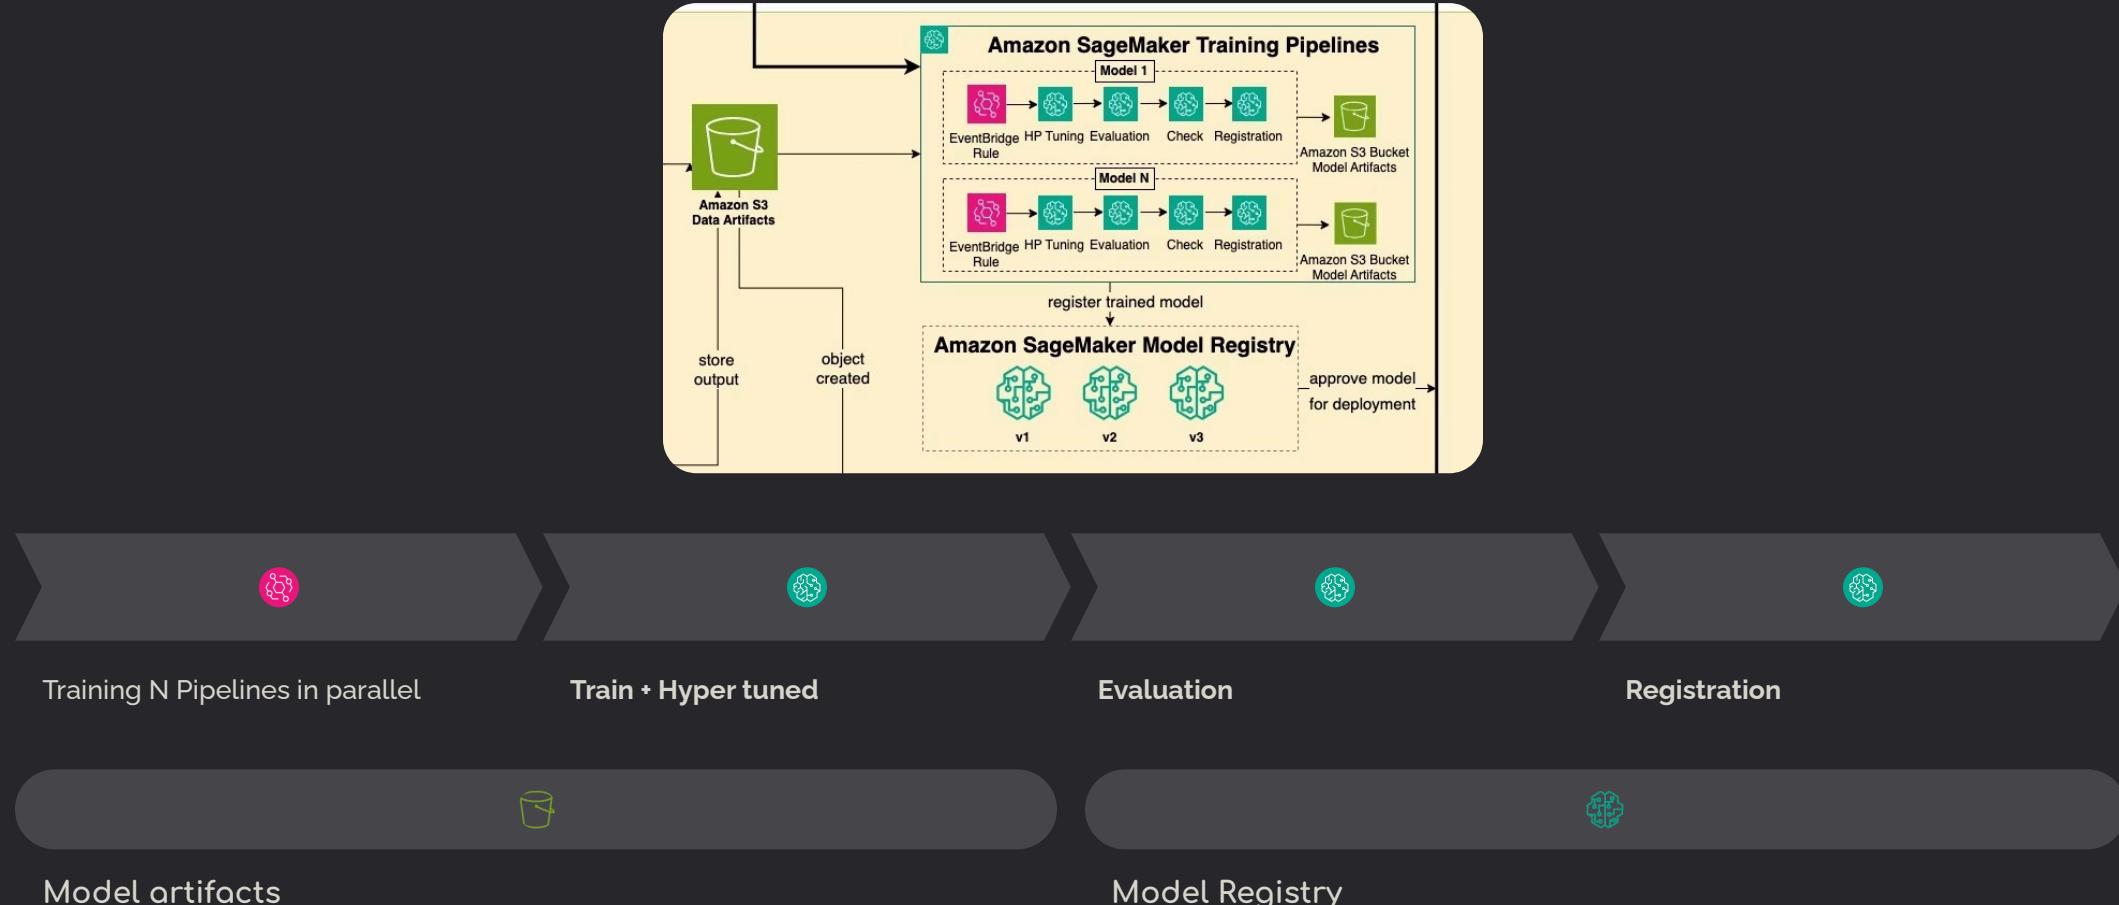
Data ingestion & Auto-Processing Pipeline

Development Account



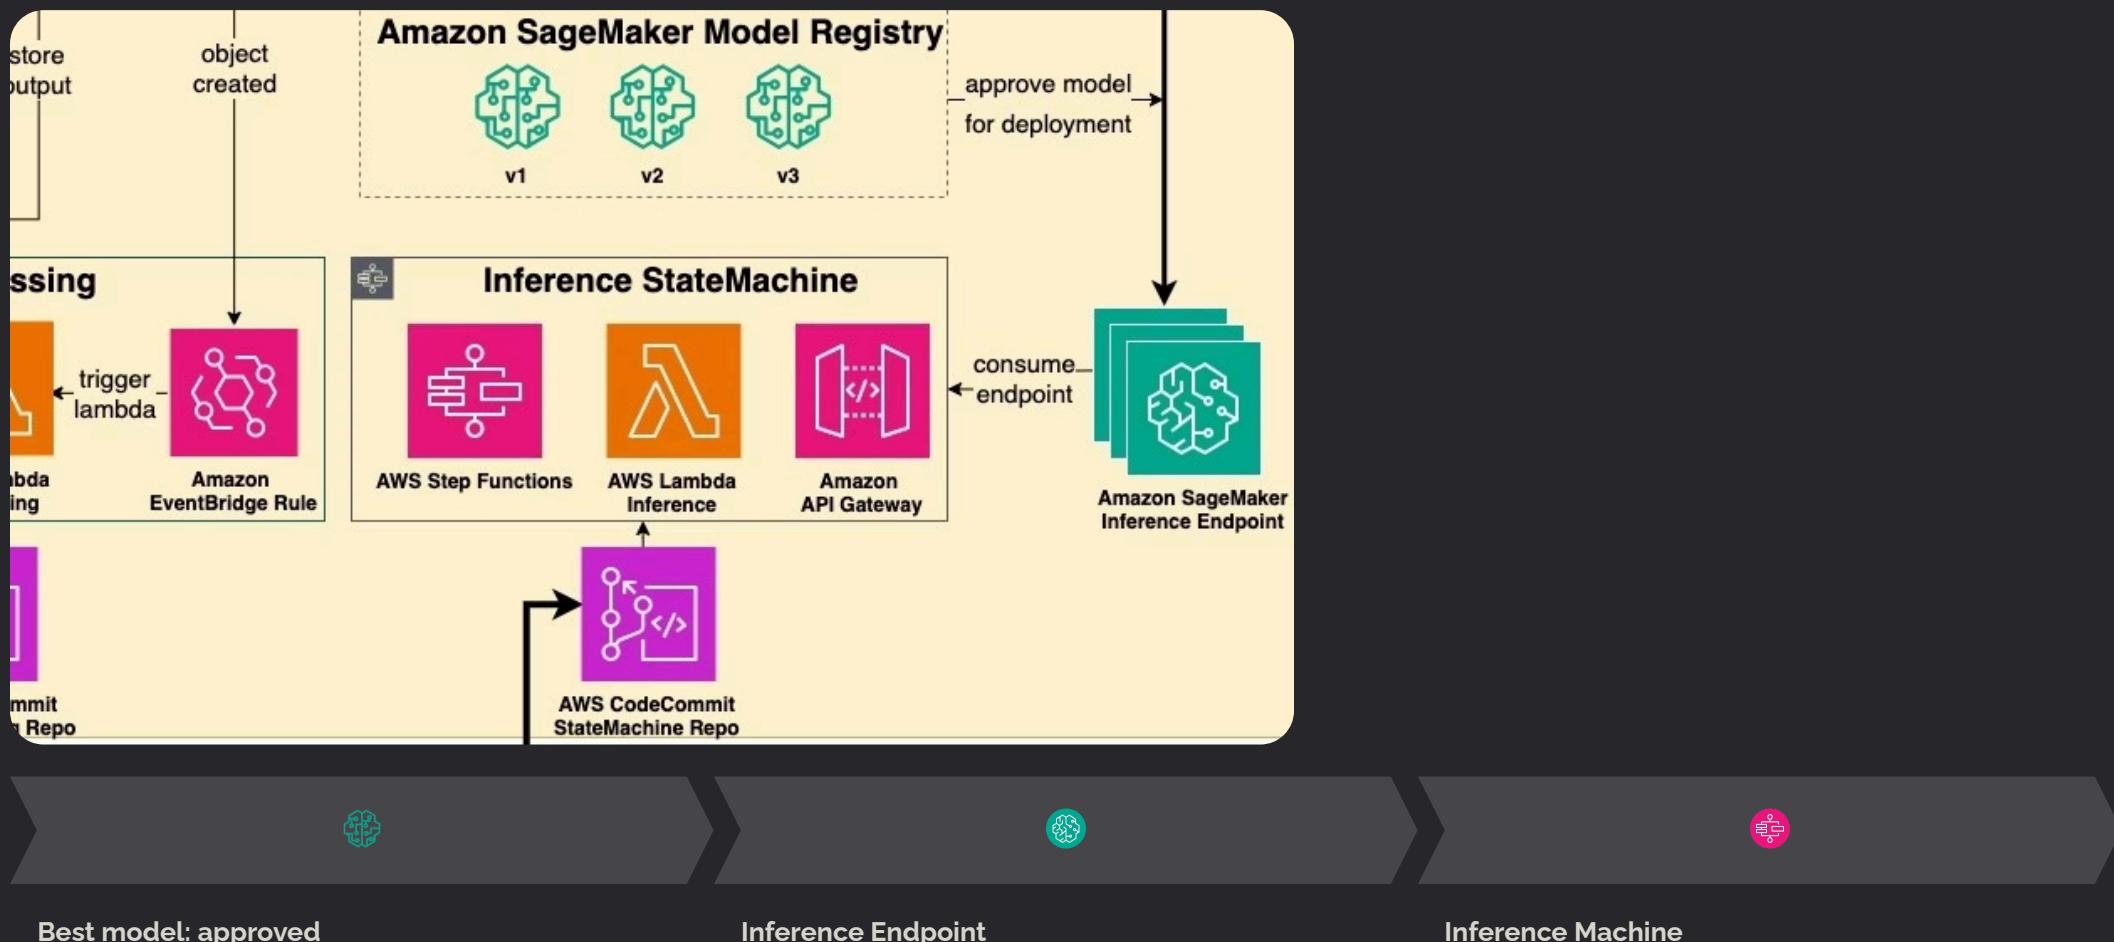
Automated Training & Tuning Pipeline

Development Account



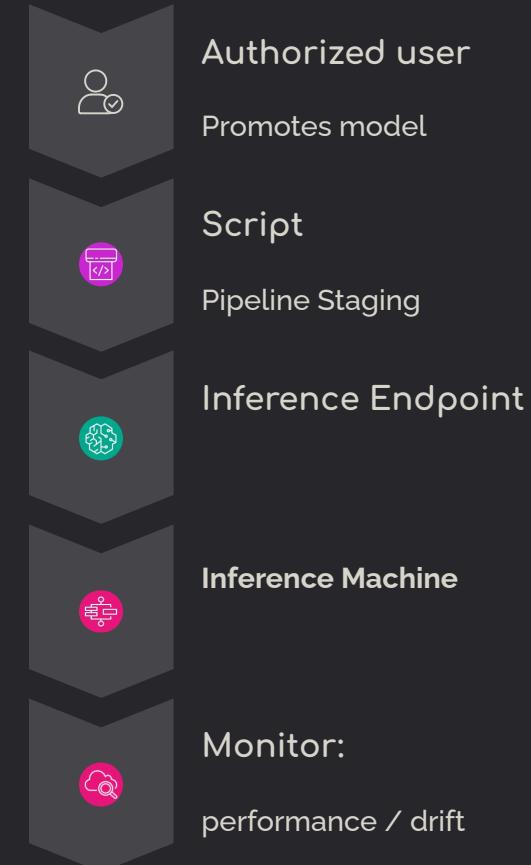
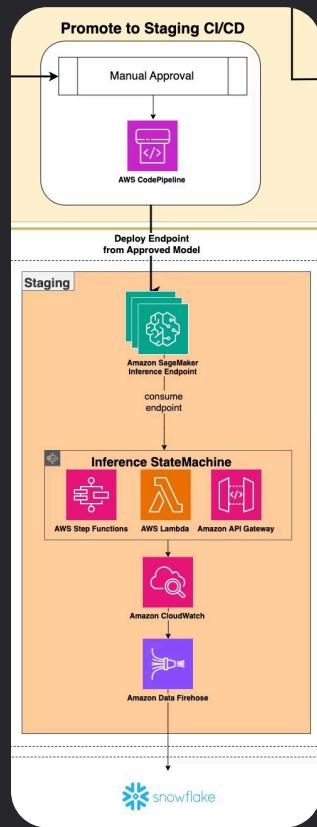
Inference Pipeline

Development Account



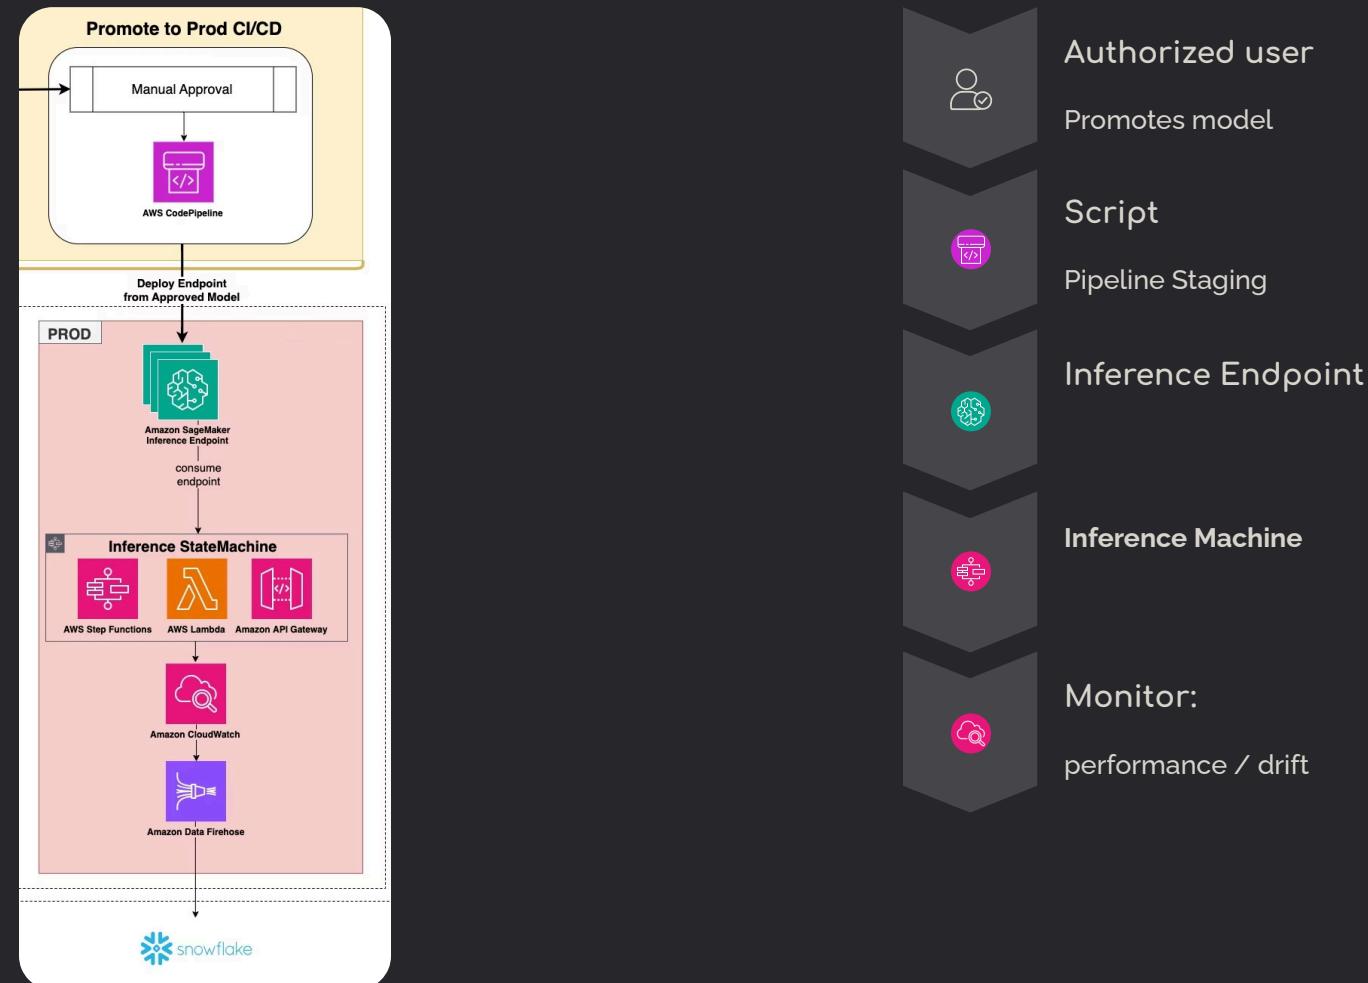
MLOps Architecture: *Staging*

Model approval and inference (i.e. beta testing)



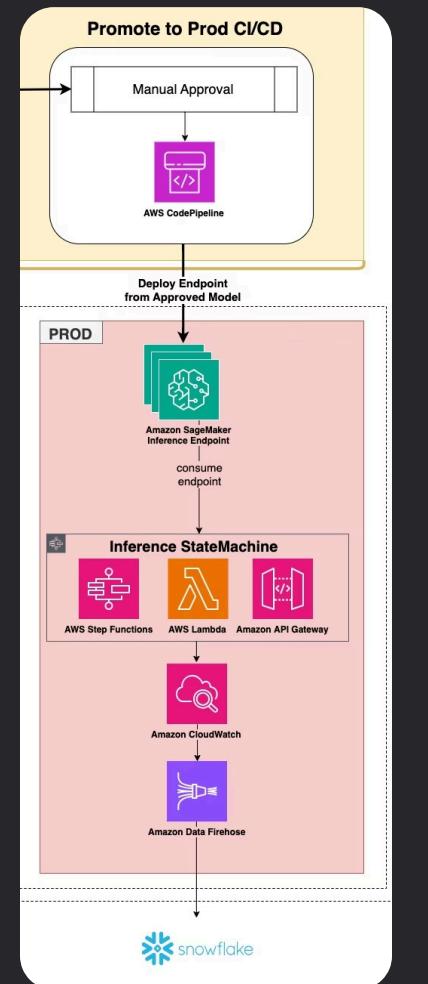
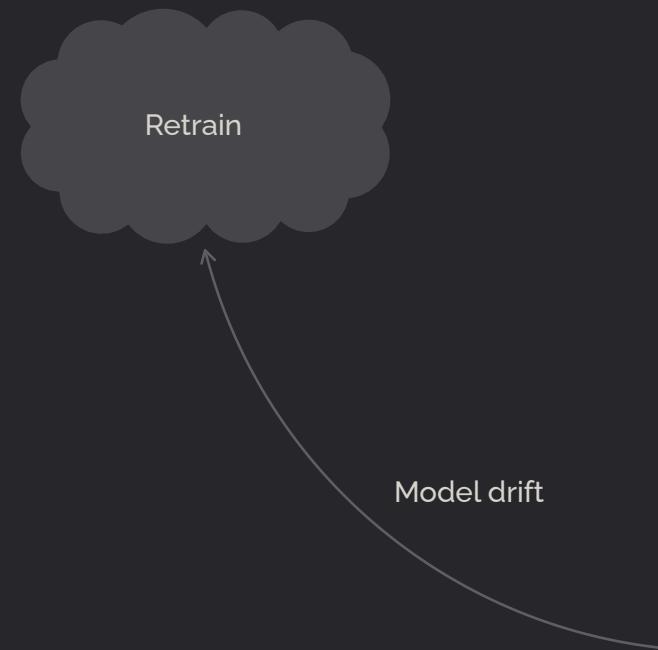
MLOps Architecture: *Production*

Model approval and inference deployment

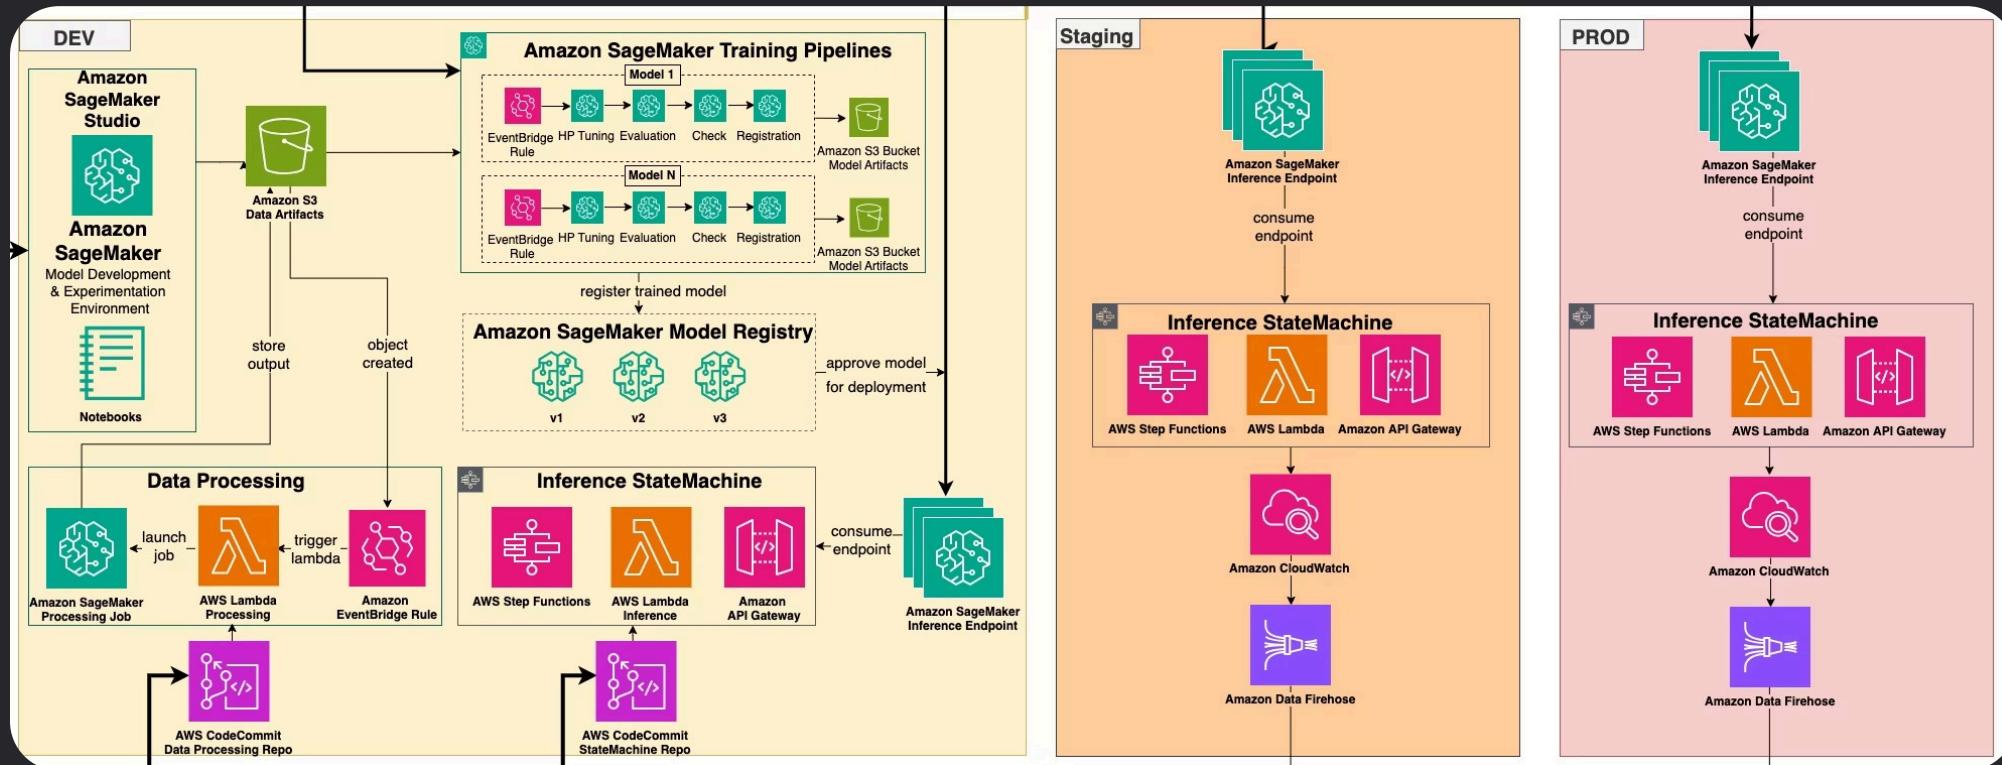


MLOps Architecture: *Production*

Model approval and inference deployment



MLOps Architecture: Automated Pipeline



Development



Staging



Production

Making AI Actionable for Claims Handlers

Car Insurance Claim



Write Off

Remove the item
from records.



SEEK REPAIR

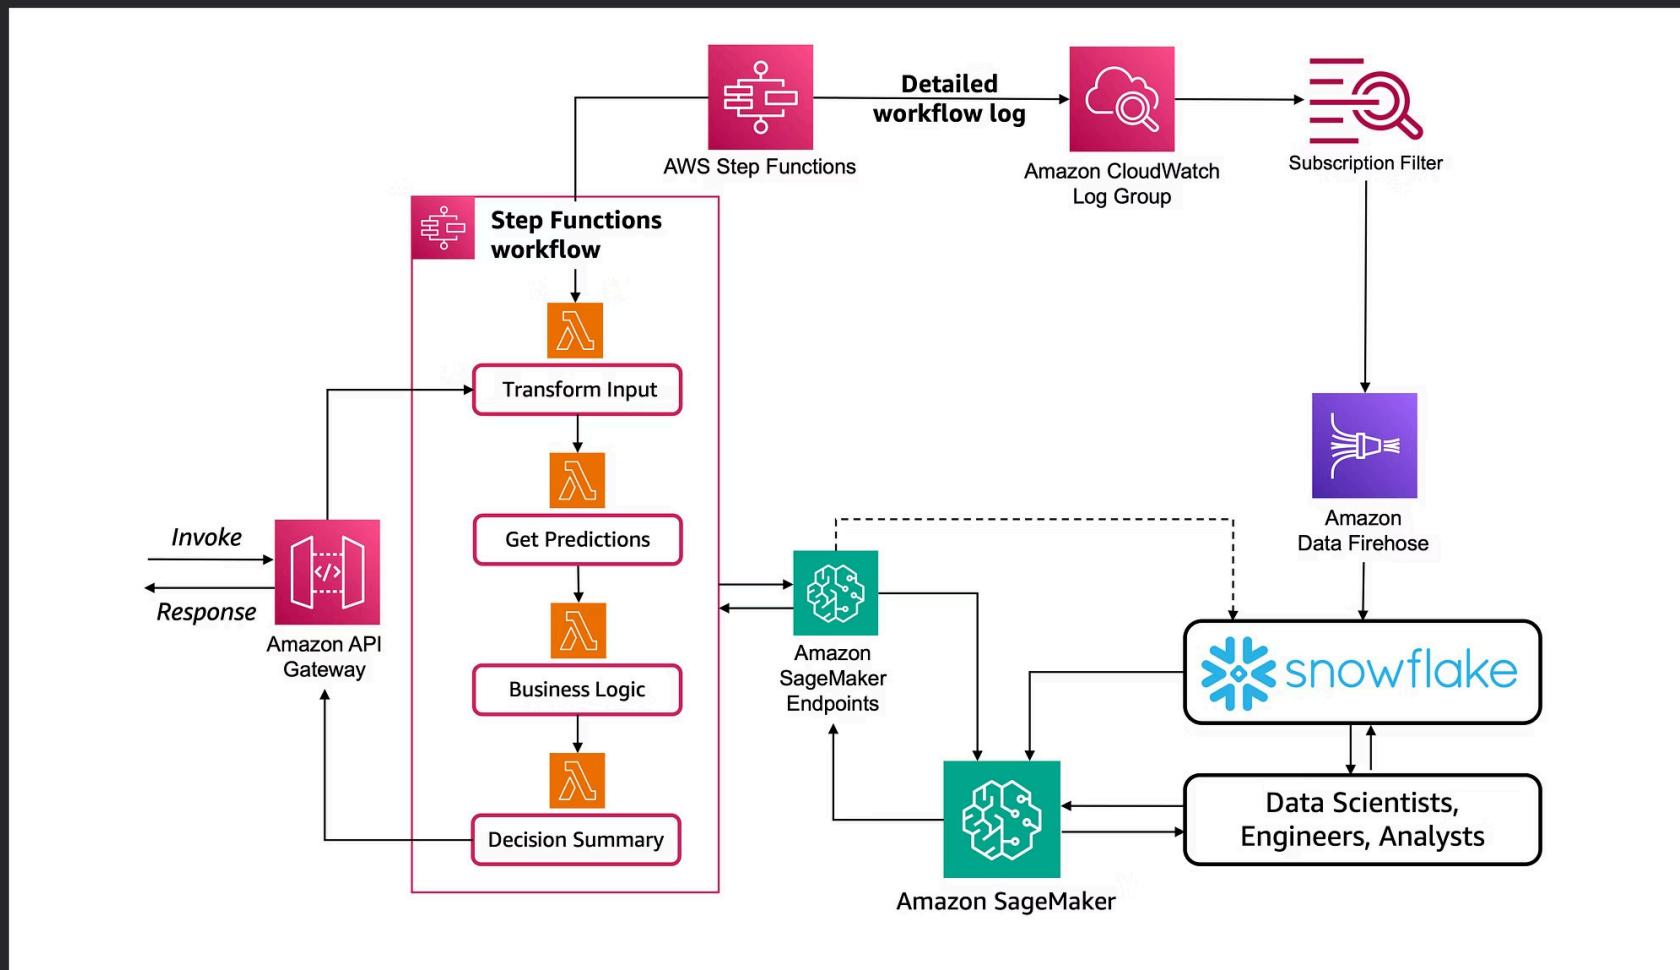
*Vehicle value >
Repair cost*

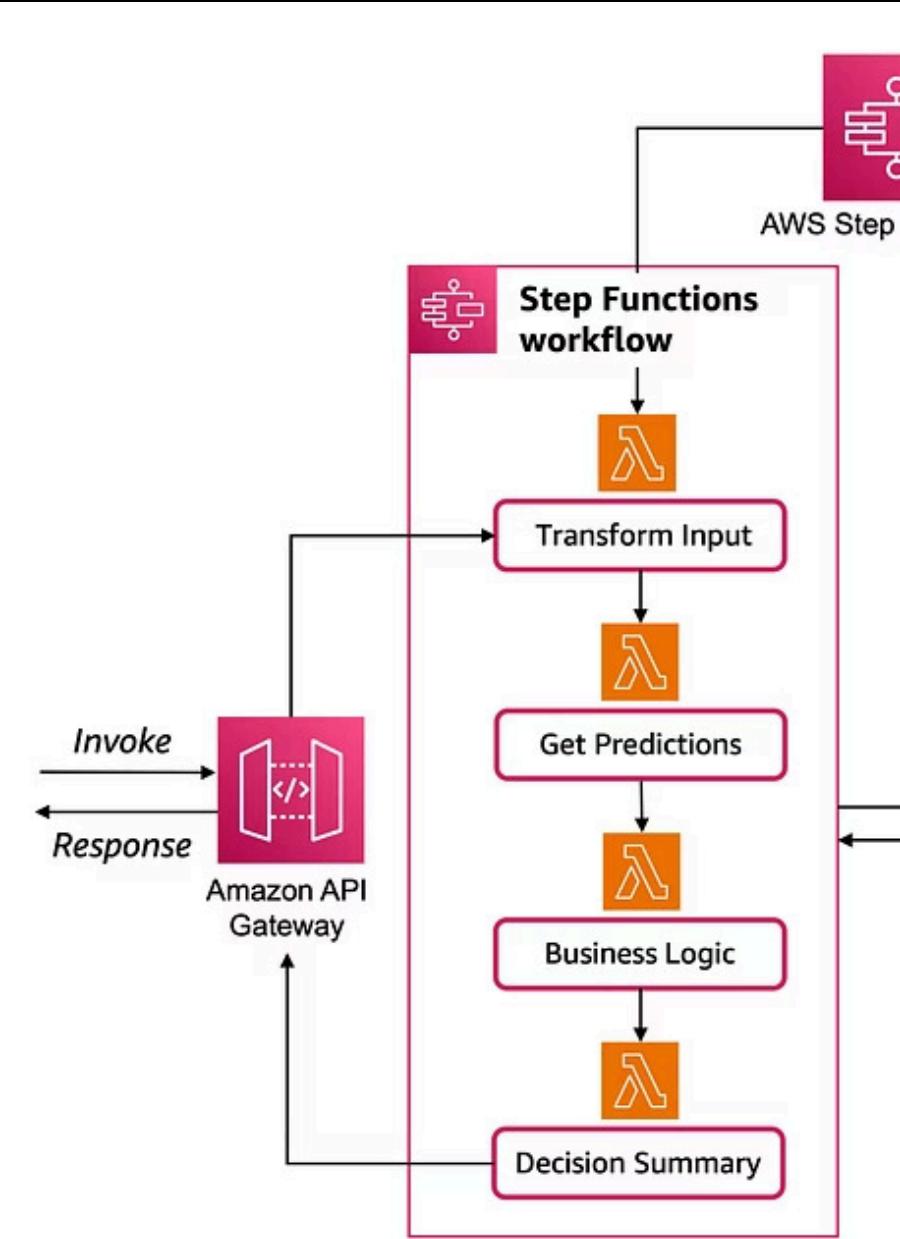


Investigate
Further

Examine to
determine the issue.

Production: Real-time Inference Pipeline



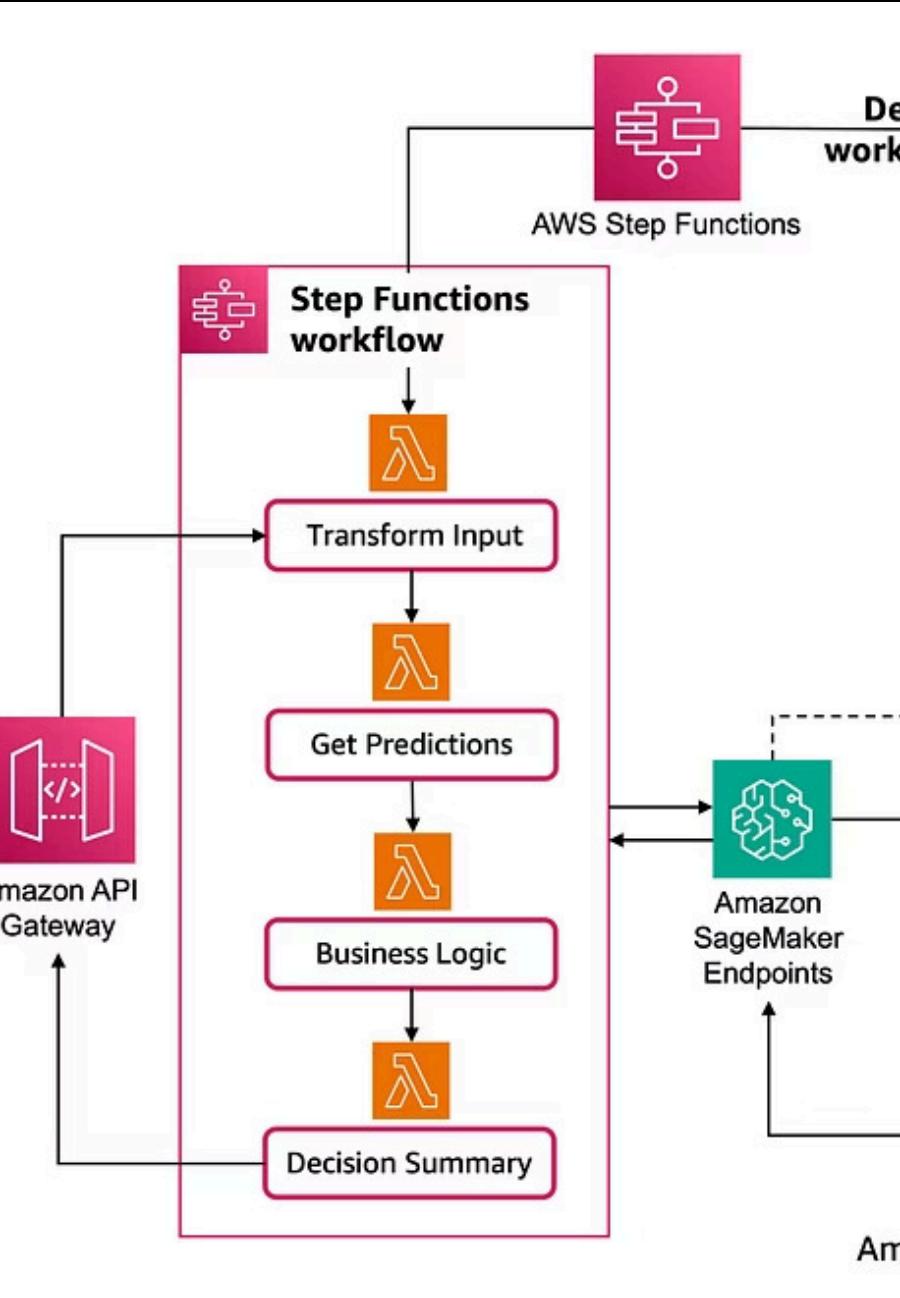


Production: Real-time Inference Pipeline



CLAIM INITIATION

Claims Handler submits Request via API



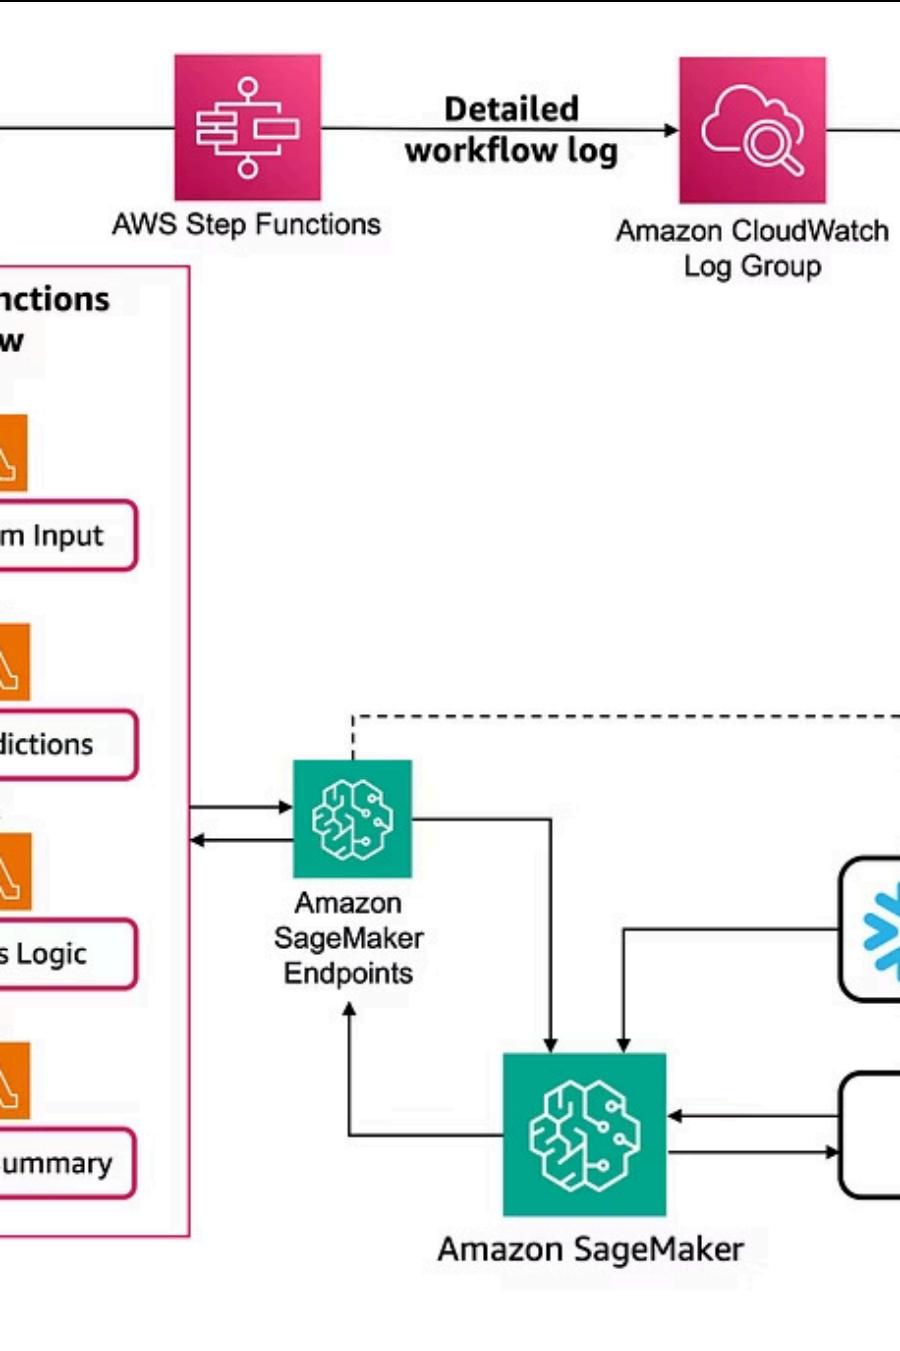
Production: Real-time Inference Pipeline

CLAIM INITIATION

Claims Handler submits Request via API

DATA TRANSFORMATION

Feature Encoding



Production: Real-time Inference Pipeline

CLAIM INITIATION

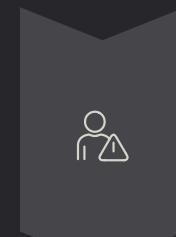
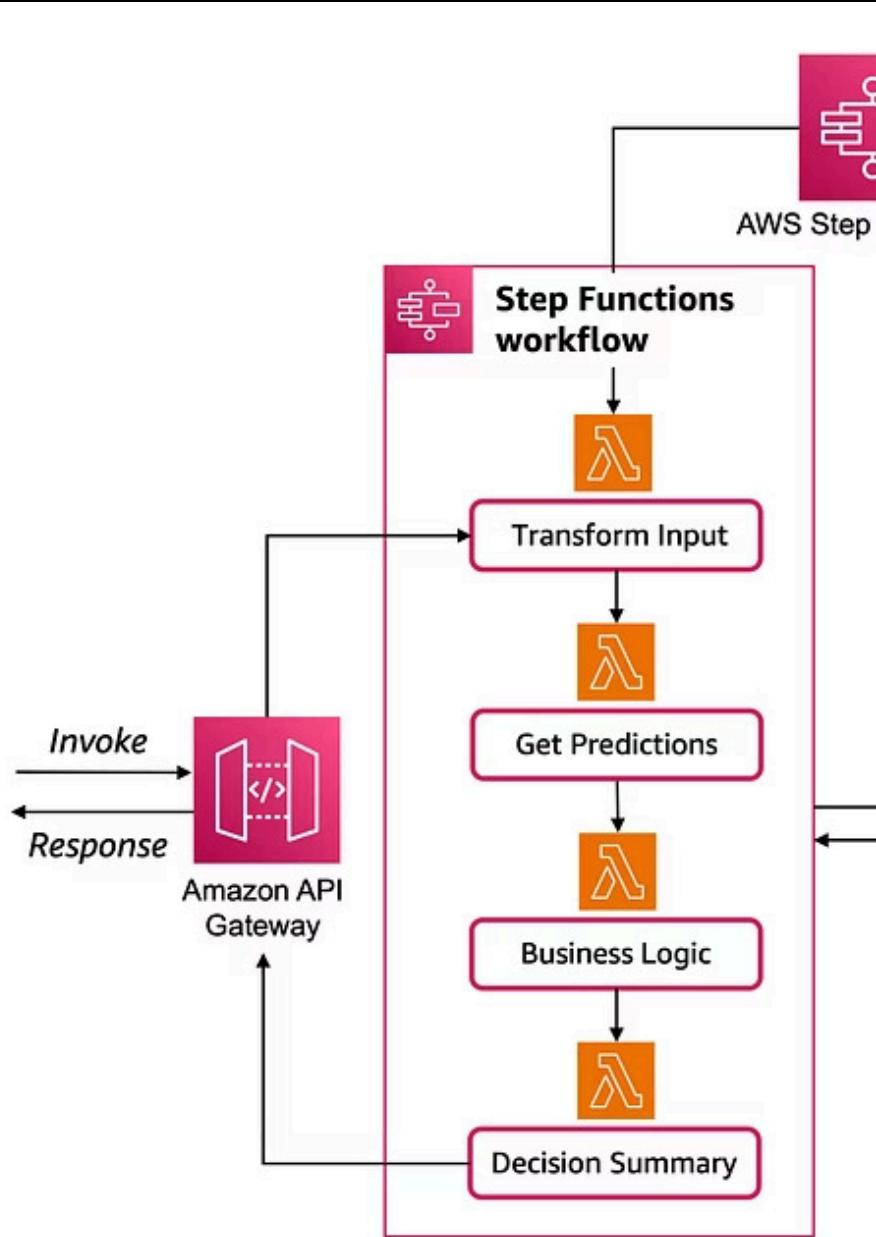
Claims Handler submits Request via API

DATA TRANSFORMATION

Feature Encoding

MODEL GENERATES PREDICTION

Production: Real-time Inference Pipeline



CLAIM INITIATION

Claims Handler submits Request via API



DATA TRANSFORMATION

Feature Encoding

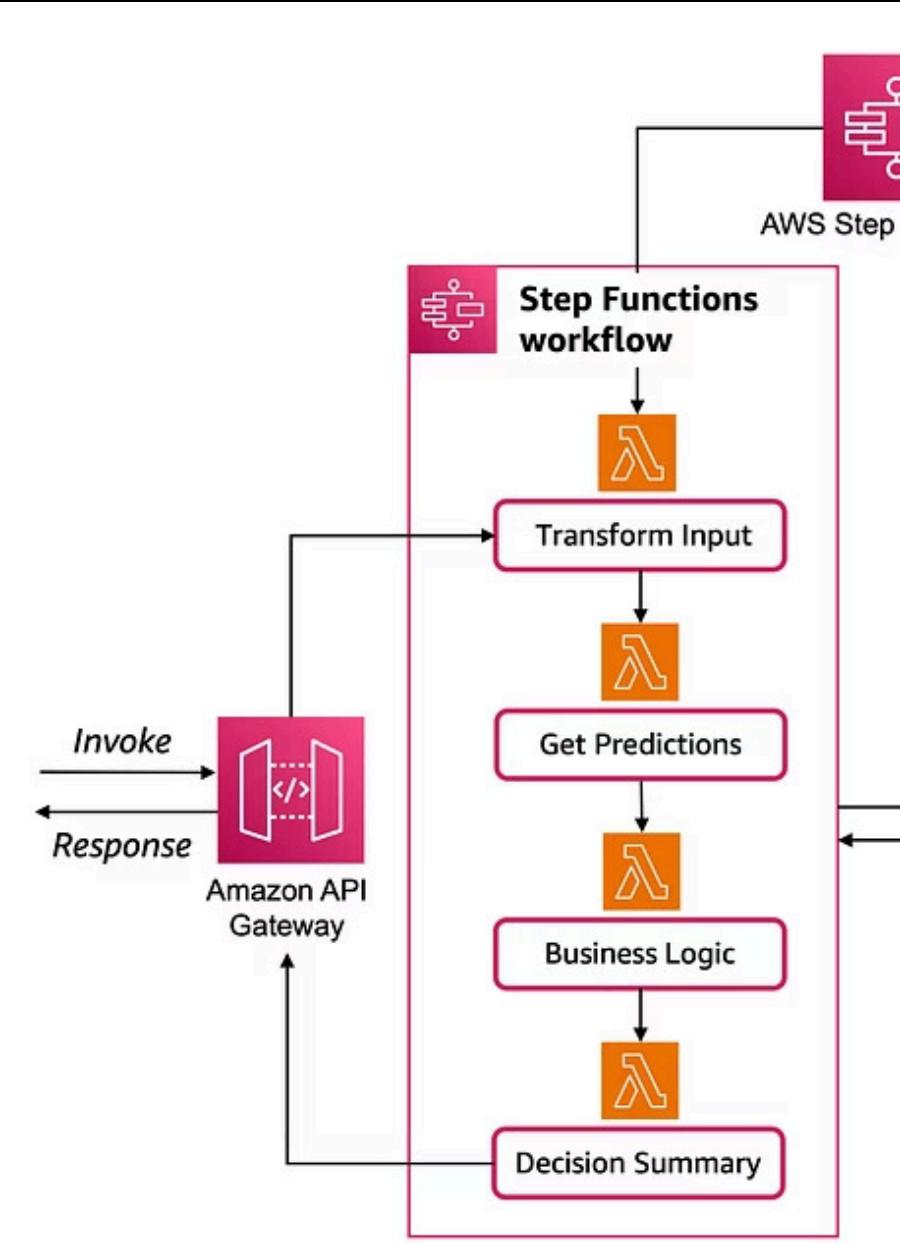


MODEL GENERATES PREDICTION

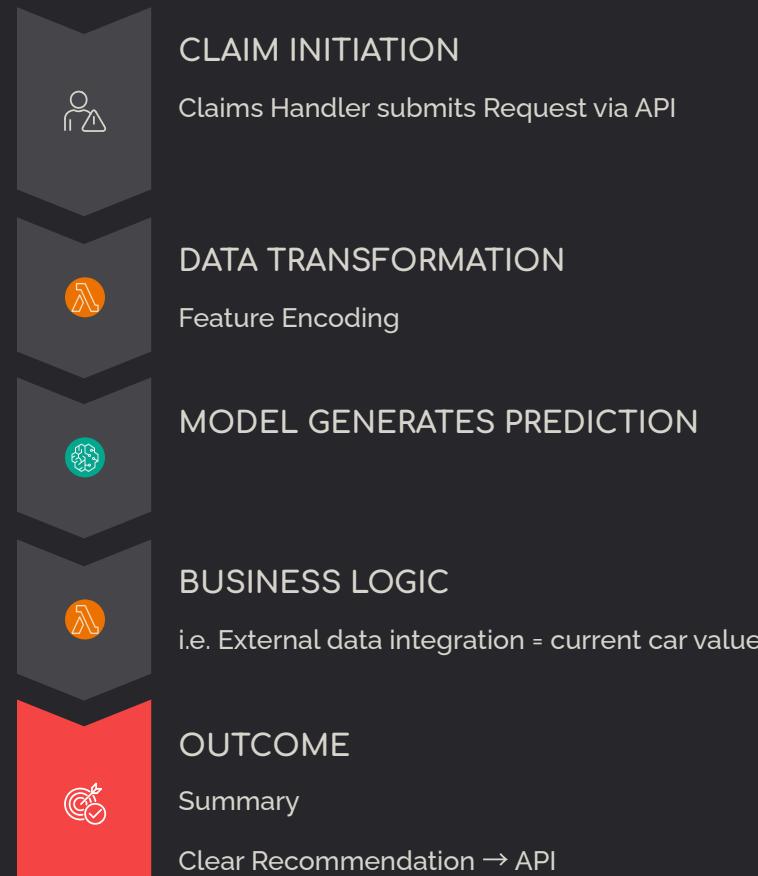


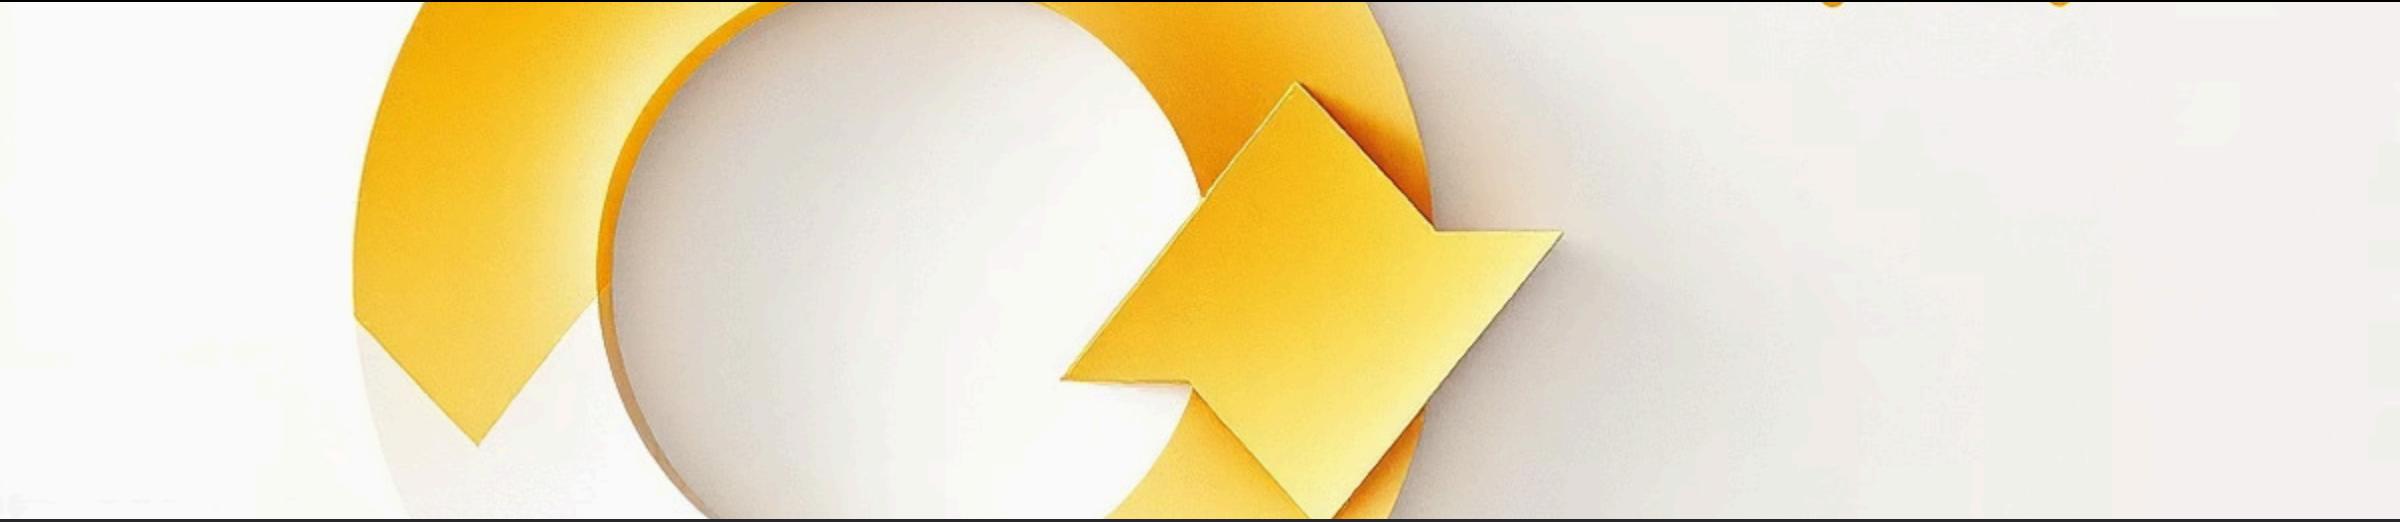
BUSINESS LOGIC

i.e. External data integration = current car value



Production: Real-time Inference Pipeline





Full Automation, Maximum Efficiency

- Pipelines Trigger Pipelines
- Human-in-the-Loop for Governance
- Continuous Monitoring Triggers Retraining



Business Impact

1

MLOps Architecture Drives Innovation

Enterprise-grade infrastructure—network isolation, unified workspaces, automated pipelines, and seed code—shifts focus from operations to innovation.

2

Cost and Speed Gains Are Real

90% infrastructure cost reduction, 100x faster deployment cycles, and 50% reduction in manual operational work demonstrate tangible business value.

3

Governance at Scale Enables Growth

Standardised templates, centralised model governance, and reproducible workflows allow Aviva to confidently **scale from 70 active use cases to hundreds across the enterprise**.

Aviva MLOps Case Study Team

Guillermo Fiallo Montero

Fee Pieper